

Assignment 2: Regression with Neural Networks

Sanaz Marefati

October 2, 2025

1 Functions and Deliverables

The functions to approximate are:

1. (1D, discontinuous) Heaviside function on $x \in [-2, 2]$.
2. (1D) $\Gamma(x) = \sqrt{2 - x}$ on $[0, 2]$, Holder continuous with exponent 1/2.
3. (2D, multi-modal) $\Gamma(x, y) = \exp(-k_1(x - 0.5)^2 - (y - 0.5)^2) + \exp(-k_2(x + 0.5)^2 - (y + 0.5)^2)$ on $[-1, 1]^2$, with $k_1 = k_2 = 1$ (optional $k_1 = k_2 = 10$).
4. Rosenbrock function $f(x, y) = 100(y - x^2)^2 + (1 - x)^2$ on $[-2, 2] \times [-2, 2]$.
5. (High-dimensions, optional) $\Gamma(x) = \frac{1}{(2\pi)^{d/2}} \exp\left(-\sum_{i=1}^d x_i^2\right)$ on $[0, 1]^d$ for $d = 1$ to 8.

Deliverables include plots of functions, NN approximations, absolute/relative errors, code, explanations of choices, and comments on best settings.

2 Methodology

Each target function was approximated using fully connected neural networks trained on quasi-Monte Carlo (QMC) Sobol points. For the 1D problems, 100 points were used, while for the 2D problems 10,000 points were generated, with independent Sobol sequences for training and testing. All inputs were normalized to $[-1, 1]$, and for the Rosenbrock function the outputs were additionally mean–std normalized to improve numerical stability.

Two architectures were employed: a shallow network with one hidden layer of 40 neurons, and a deep network with three hidden layers of 20 neurons each (comparable total neuron count). Both `ReLU` and `tanh` activations were tested, using He uniform and variance scaling initialization respectively. Training was performed with the ADAM optimizer and a polynomially decaying learning rate ($10^{-3} \rightarrow 10^{-4}$) over 10,000 epochs. A summary of these settings is provided in Table 1.

Remark. For the Rosenbrock function, the initial error was large due to insufficient output normalization. Applying mean–std scaling to the outputs prior to training, and denormalizing for error computation, significantly improved model performance.

Table 1: Summary of neural network settings and best performers for each function.

Function	Heaviside	$\sqrt{2-x}$	2D Gaussians	Rosenbrock
Training points	100 (QMC Sobol)	100 (QMC Sobol)	10,000 (QMC Sobol)	10,000 (QMC Sobol)
Test points	100	100	10,000	10,000
NN architecture	Shallow (1,40,1)	Deep (1,20,20,20,1)	Deep (2,20,20,20,1)	Deep (2,20,20,20,1)
Activation function	ReLU	tanh	ReLU	ReLU
Initialization	He_uniform (ReLU)	Variance Scaling (tanh)	Variance Scaling (tanh)	He_uniform (ReLU)
Optimizer	ADAM (lr=10 ⁻³ , decay to 10 ⁻⁴)	ADAM (lr=10 ⁻³ , decay to 10 ⁻⁴)	ADAM (lr=10 ⁻³ , decay to 10 ⁻⁴)	ADAM (lr=10 ⁻³ , decay to 10 ⁻⁴)
Epochs	10,000	10,000	10,000	10,000
Best performer	Shallow ReLU	Deep tanh	Deep tanh	Deep ReLU

2.1 Code for the Problems

```
1 #===== Setup =====
2 import jax, jax.numpy as jnp
3 from jax import random, vmap, value_and_grad, jit, grad, nn,
4     jacrev
5 from jax.nn import relu, tanh, initializers
6 from jax.nn.initializers import variance_scaling, he_normal,
7     he_uniform
8 import optax
9 from scipy.stats import qmc
10 from scipy.interpolate import griddata
11 import matplotlib.pyplot as plt
12 from matplotlib.colors import Normalize, LightSource
13 from mpl_toolkits.mplot3d import Axes3D
14 import time
15
16 #===== Define functions and datasets =====
17 def f_heaviside(x):
18     return jnp.where((x >= 0.0) & (x <= 2.0), 1.0, 0.0)
19
20 def f_gamma(x):
21     return jnp.sqrt(jnp.maximum(2.0 - x, 0.0))
22
23 def f_gaussians(x1, x2):
24     k1, k2 = 1.0, 1.0
25     #k1, k2 = 10, 10
26     return jnp.exp(-(k1*(x1 - 0.5)**2 + (x2 - 0.5)**2)) + \
27             jnp.exp(-(k2*(x1 + 0.5)**2 + (x2 + 0.5)**2))
28
29 def f_rosenbrock(x1, x2):
30     return 100*(x2 - x1**2)**2 + (1-x1)**2
31
32 def f_gaussian_highdim(x):
33     d = x.shape[-1]
34     prefactor = 1 / (2 * jnp.pi)**(d / 2)
35     return prefactor * jnp.exp(-jnp.sum(x**2, axis=-1))
36
37 datasets = {
38     'heaviside': {
39         'dim': 1,
40         'domain': [(-2, 2)],
41         'func': f_heaviside,
42         'name': 'Heaviside'
43     },
44     'gamma': {
45         'dim': 1,
46         'domain': [(0, 2)],
47         'func': f_gamma,
48         'name': 'sqrt(2-x)'
49     },
50 }
```

```

48     'gaussians': {
49         'dim': 2,
50         'domain': [(-1, 1), (-1, 1)],
51         'func': f_gaussians,
52         'name': '2D Gaussians'
53     },
54     'rosenbrock': {
55         'dim': 2,
56         'domain': [(-2, 2), (-2, 2)],
57         'func': f_rosenbrock,
58         'name': 'Rosenbrock'
59     }
60 }
61 #high_dim_xs = lambda d: {'f': f_gaussian_highdim, 'dim': d,
62 #    'domain': [(0, 1)]*d, 'name': f'High-d Gaussian (d={d})'}
63
64 def bounds(domain):
65     lower = jnp.array([d[0] for d in domain])
66     upper = jnp.array([d[1] for d in domain])
67     return lower, upper
68
69 def generate_qmc_samples(dim, n, domain, seed=0):
70     sampler = qmc.Sobol(d=dim, scramble=True, seed=seed)
71     samples = sampler.random(n=n)
72     lower, upper = bounds(domain)
73     samples = lower + samples * (upper - lower)
74     return jnp.array(samples)
75
76 def normalize(x, domain):
77     lower, upper = bounds(domain)
78     return 2 * (x - lower) / (upper - lower + 1e-10) - 1
79
80 # ===== Neural Network Architecture =====
81 def initialize_mlp(sizes, key, activ_name):
82     keys = random.split(key, len(sizes))
83
84     def initialize_layer(m, n, k, activ_name):
85         w_key, b_key = random.split(k)
86         if activ_name == 'relu':
87             initializer = he_uniform() # He initialization for
88             #relu
89             #initializer = he_normal()
90         elif activ_name == 'tanh':
91             initializer = variance_scaling(1.0, 'fan_avg',
92             'uniform')
93             #initializer = glorot_uniform()
94         else:
95             raise ValueError("Unsupported activation.")
96     W = initializer(w_key, (m, n))
97     b = jnp.zeros((n,))
98     return W, b

```

```

96
97     return [initialize_layer(m, n, k, activ_name) for m, n, k in
98             zip(sizes[:-1], sizes[1:], keys)]
99
100    # Forward pass
101   def forward(params, acts, activation):
102       for w, b in params[:-1]:
103           acts = activation(jnp.dot(acts, w) + b)
104       w, b = params[-1]
105       return jnp.dot(acts, w) + b
106
107    # Predict function
108   def predict(params, x, activation):
109       pred = vmap(lambda x_i: forward(params, x_i, activation))(x)
110       return jnp.squeeze(pred, axis=-1) # Shape: (n_data,)
111
112    # Loss function
113   def loss(params, x, y, activation):
114       y_pred = predict(params, x, activation)
115       return jnp.mean(jnp.square(y_pred - y))
116
117    # ===== Training with optax (ADAM with scheduler) =====
118   def train_model(x_train, y_train, sizes, activation,
119                   epochs=10000, lr=1e-3, batch_size=32, seed=0):
120       key = random.PRNGKey(seed)
121       params = initialize_mlp(sizes, key, activation)
122
123       schedule = optax.polynomial_schedule(init_value=lr,
124                                             end_value=lr * 0.1, power=2.0, transition_steps=1000 )
125       # Set up optimizer in Optax
126       optimizer = optax.adam(learning_rate=schedule)
127       opt_state = optimizer.init(params)
128
129       loss_grad = value_and_grad(loss, argnums=0)
130       loss_history = []
131
132       n_batches = len(x_train) // batch_size
133       for epoch in range(epochs):
134           epoch_loss = 0.0
135           perm = random.permutation(key, len(x_train))
136           x_train = x_train[perm]
137           y_train = y_train[perm]
138           for i in range(n_batches):
139               start = i * batch_size
140               end = start + batch_size
141               x_batch = x_train[start:end]
142               y_batch = y_train[start:end]
143               loss_val, grads = loss_grad(params, x_batch,
144                                           y_batch, activations[activ_name])
145               updates, opt_state = optimizer.update(grads,
146                                                     opt_state)

```

```

142     params = optax.apply_updates(params, updates)
143     epoch_loss += loss_val
144     loss_history.append(epoch_loss / n_batches)
145     if epoch % 1000 == 0:
146         print(f'Epoch [{epoch}/{epochs}], Loss: {loss_history[-1]:.6f}')
147     return params, loss_history
148
149 # ===== Main Execution =====
150 start_time = time.time() # Start timing
151
152 # Choose dataset
153 #selected_dataset = datasets['heaviside']
154 #selected_dataset = datasets['gamma']
155 selected_dataset = datasets['gaussians']
156 #selected_dataset = datasets['rosenbrock']
157
158 dim = selected_dataset['dim']
159 domain = selected_dataset['domain']
160 func = selected_dataset['func']
161 name = selected_dataset['name']
162
163 data_points=100 if dim == 1 else 10000
164 test_points=100 if dim == 1 else 10000
165
166 x_data = generate_qmc_samples(dim, data_points, domain, seed=0)
167 x_test = generate_qmc_samples(dim, test_points, domain, seed=1)
168 x_test_norm = normalize(x_test, domain)
169
170 y_data = func(x_data[:, 0]) if dim == 1 else
171     func(*x_data.T[:dim])
172 y_test = func(x_test[:, 0]) if dim == 1 else
173     func(*x_test.T[:dim])
174
175 arch_shallow = [dim, 40, 1]           # 60 neurons total in hidden
176 arch_deep   = [dim, 20, 20, 20, 1] # 20*3 = 60 neurons total
177 for activ_name in ['relu', 'tanh']:
178     for name, size in [('shallow', arch_shallow), ('deep',
179                         arch_deep)]:
180         params, hist = train_model(x_test_norm, y_test, size,
181                                     activ_name, epochs=10000, lr=1e-3)
182         y_pred = predict(params, x_test_norm,
183                           activations[activ_name])
184
185 abs_err = jnp.abs(y_pred - y_test)
186 rel_err = abs_err / (jnp.abs(y_test) + 1e-10)
187
188 print(f'{name} - Mean Absolute Error: {jnp.mean(abs_err):.6f},
189       Mean Relative Error: {jnp.mean(rel_err):.6f}')
190
191 end_time = time.time()

```

```

186 print(f"Total Execution time: {end_time - start_time} seconds")
187
188 # ===== Visualization =====
189 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
190 axes = axes.flatten()
191 plt.rcParams.update({'font.size': 10})
192
193 if dim == 1:
194     axes[0].scatter(x_data[:, 0], y_data, c='b', s=10,
195                      alpha=0.8, label='True')
196     axes[0].set_xlabel('Data Points')
197     axes[0].set_ylabel('True Function')
198     axes[0].grid(True, alpha=0.3)
199
200     axes[1].scatter(x_test[:, 0], y_pred, c='orange', s=10,
201                      alpha=0.8, label='NN Approx')
202     axes[1].set_xlabel('Data Points')
203     axes[1].set_ylabel('NN Approximation')
204     axes[1].grid(True, alpha=0.3)
205
206     axes[2].scatter(x_test[:, 0], abs_err, c='r', s=10,
207                      alpha=0.8)
208     axes[2].set_xlabel('Data Points')
209     axes[2].set_ylabel('Absolute Error')
210     axes[2].grid(True, alpha=0.3)
211
212     axes[3].scatter(x_test[:, 0], rel_err, c='g', s=10,
213                      alpha=0.8)
214     axes[3].set_xlabel('Data Points')
215     axes[3].set_ylabel('Relative Error')
216     axes[3].grid(True, alpha=0.3)
217
218 else:
219     x1, x2 = x_data[:, 0], x_data[:, 1]
220     x1i, x2i = jnp.linspace(min(x1), max(x1), 100),
221                 jnp.linspace(min(x2), max(x2), 100)
222     xi, xii = jnp.meshgrid(x1i, x2i)
223     points = jnp.stack([xi.ravel(), xii.ravel()], axis=1)
224
225     y_true_grid = func(points[:, 0], points[:, 1]).reshape(xi.shape) #griddata((x1, x2), y_data,
226     (xi, xii), method='cubic')
227     y_pred_grid = predict(params, points,
228                           activations[activ_name]).reshape(xi.shape)
229     #((x_test[:, 0], x_test[:, 1]), y_pred, (xi, xii),
230     method='cubic')
231
232     abs_err_grid = jnp.abs(y_pred_grid - y_true_grid)
233     rel_err_grid = abs_err_grid / (jnp.abs(y_true_grid) +
234                                   1e-10)
235
236     cmap_true = 'viridis' if 'Gaussian' in name else 'plasma'

```

```

226
227     axes[0].scatter(x1, x2, c=y_data, cmap='viridis' if name
228         != 'Rosenbrock' else 'plasma', s=5)
229     axes[0].set_xlabel('x1')
230     axes[0].set_ylabel('x2')
231     cbar = plt.colorbar(axes[0].collections[0], ax=axes[0])
232     cbar.set_label('True Function', rotation=90)
233
234     cont = axes[1].contourf(xi, xii, y_pred_grid,
235         levels=100, cmap='viridis' if name != 'Rosenbrock'
236         else 'plasma')
237     axes[1].set_xlabel('x1')
238     axes[1].set_ylabel('x2')
239     cbar = plt.colorbar(cont, ax=axes[1])
240     cbar.set_label('NN Approximation', rotation=90)
241
242     axes[2].remove()
243     axes[2] = fig.add_subplot(2, 2, 3, projection='3d')
244     axes[2].plot_surface(xi, xii, abs_err_grid,
245         cmap='Reds', linewidth=0, antialiased=True)
246     axes[2].set_xlabel('x1')
247     axes[2].set_ylabel('x2')
248     axes[2].set_zlabel('Absolute Error')
249     #axes[2].grid(True, alpha=0.3)
250
251     axes[3].remove()
252     axes[3] = fig.add_subplot(2, 2, 4, projection='3d')
253     axes[3].plot_surface(xi, xii, rel_err_grid,
254         cmap='Greens', linewidth=0, antialiased=True)
255     axes[3].set_xlabel('x1')
256     axes[3].set_ylabel('x2')
257     axes[3].set_zlabel('Relative Error')
258     axes[3].grid(True, alpha=0.3)
259
260 fig.suptitle(f'{name} | Activation: {choice_act} | Layers NN:
261   {choice_arch} | Training pts: {test_points}')
262 plt.tight_layout()
263 plt.savefig(f'{name} | Activation: {choice_act} | Layers NN:
264   {choice_arch} | Training pts: {test_points}', dpi=300)
265 plt.show()

```

3 Results and Comments

Plots of the functions and their neural network approximations are shown in figures. Each figure includes the true function, the NN prediction, and the absolute and relative error. Networks were trained with QMC Sobol points, using the ADAM optimizer with a polynomially decaying learning rate ($10^{-3} \rightarrow 10^{-4}$), and both ReLU and tanh activations. Two architectures were considered: a shallow model with one hidden layer (40 units) and a deep model with three hidden layers (20 units each).

3.1 Function-wise Performance

- **Heaviside (1D)**: Shallow `ReLU` performed best, reflecting the piecewise linear nature of the function.
- $\sqrt{2-x}$ (1D): Deep `tanh` networks achieve lower overall error due to smooth approximation, but at $x = 2$ `ReLU` yields smaller local error, showing its advantage in handling sharp slope changes.
- **2D Gaussians**: Deep `tanh` captured the multi-modal structure most effectively, both for $k = 1$ and $k = 10$.
- **Rosenbrock (2D)**: Deep `ReLU` performed best, handling the narrow curved valley after output normalization.

Summary: Shallow `ReLU` networks are well-suited for discontinuous or piecewise functions, while deep `tanh` networks are superior for smooth, complex, and higher-dimensional problems. In the case of the 2D Gaussians, when $k = 10$ the peaks become narrower and sharper compared to $k = 1$, creating highly localized structures. This makes the approximation problem more challenging and increases the need for deeper architectures and smooth activations (`tanh`). Shallow or `ReLU`-based networks tend to underfit in this setting, whereas deep `tanh` networks achieve the lowest errors.

For the Rosenbrock function, `ReLU` sometimes converges faster under a limited training budget. However, with longer training and improved normalization, `tanh` networks generally catch up and surpass `ReLU` in capturing the nonlinear valley.

3.2 Comments on Performance

- `ReLU` converges quickly but is less accurate for smooth functions.
- `tanh` converges more slowly but achieves lower final errors on continuous problems.
- Deep networks outperform shallow ones for multi-modal or nonlinear landscapes.
- QMC Sobol sampling improves coverage and stability compared to random sampling.

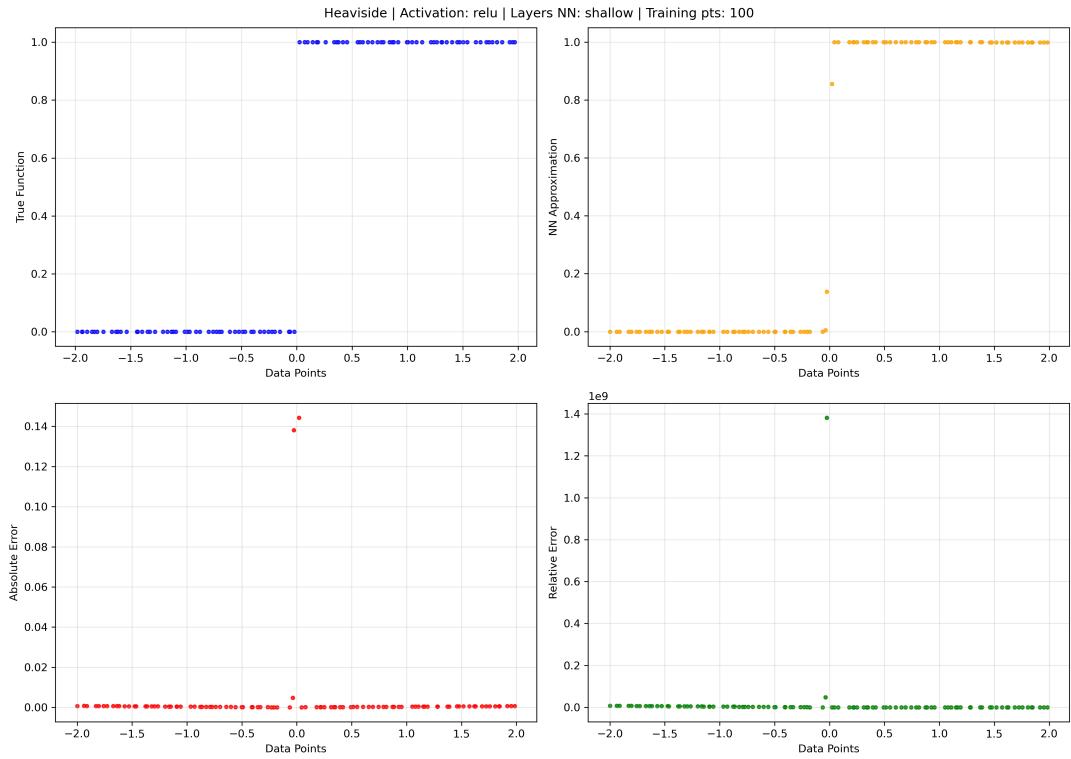


Figure 1: Heaviside step function results for shallow ReLU which performs best and reflects its suitability for piecewise functions.

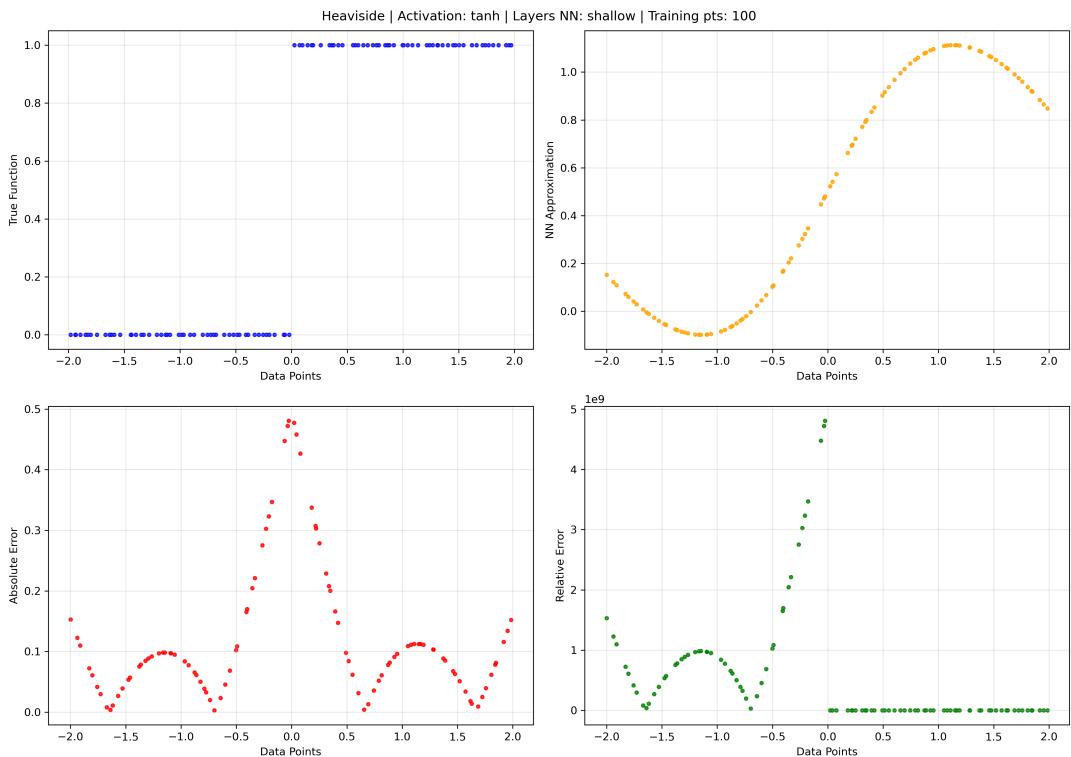


Figure 2: Heaviside step function results for Shallow tanh.

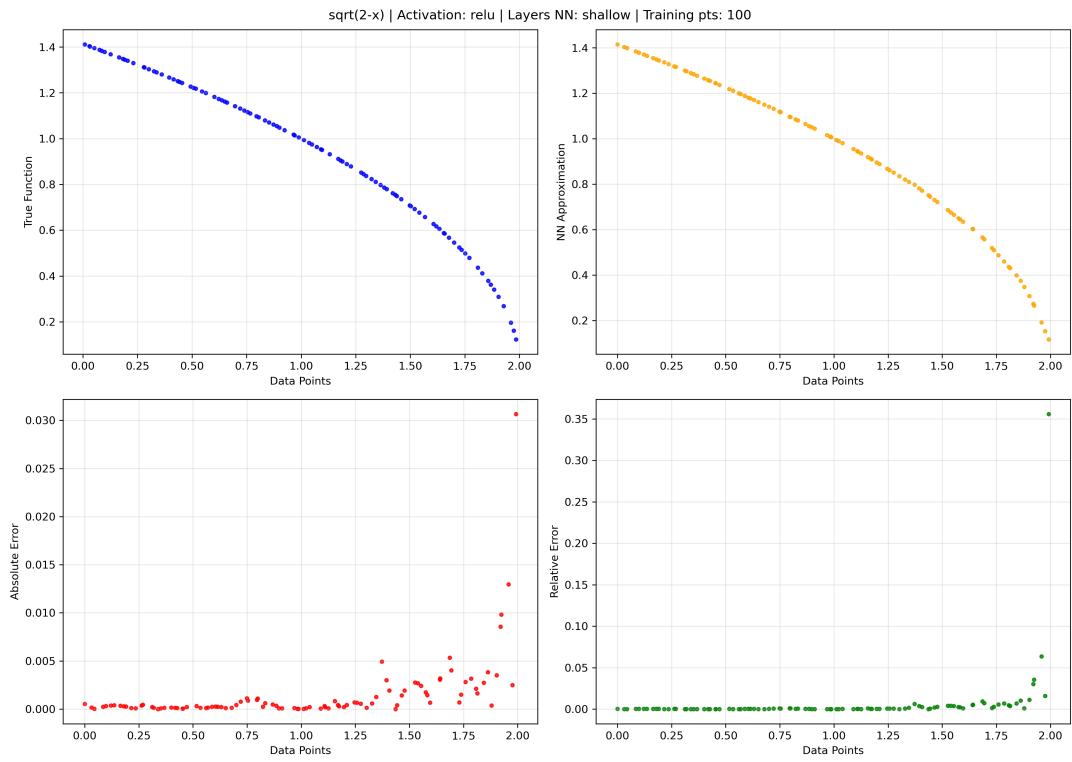


Figure 3: $\sqrt{(2 - x)}$ function results for shallow ReLU.

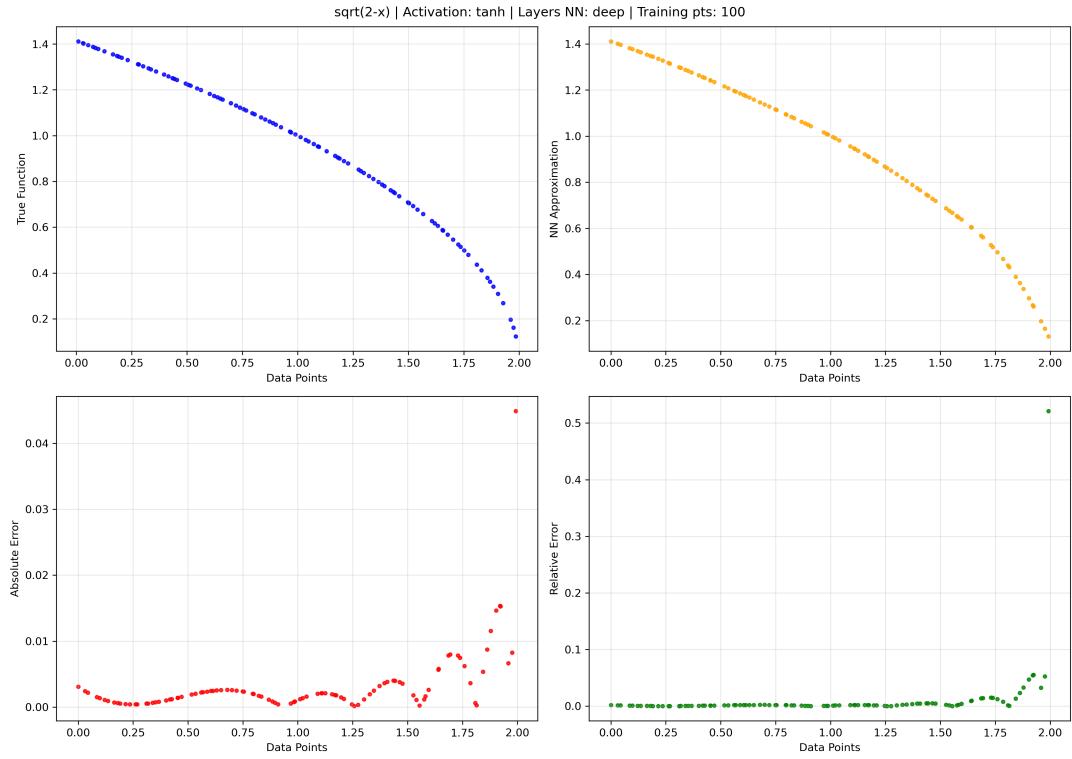


Figure 4: $\sqrt{(2 - x)}$ function results for deep tanh which performs best and reflects its smoothness for this Holder continuous function.

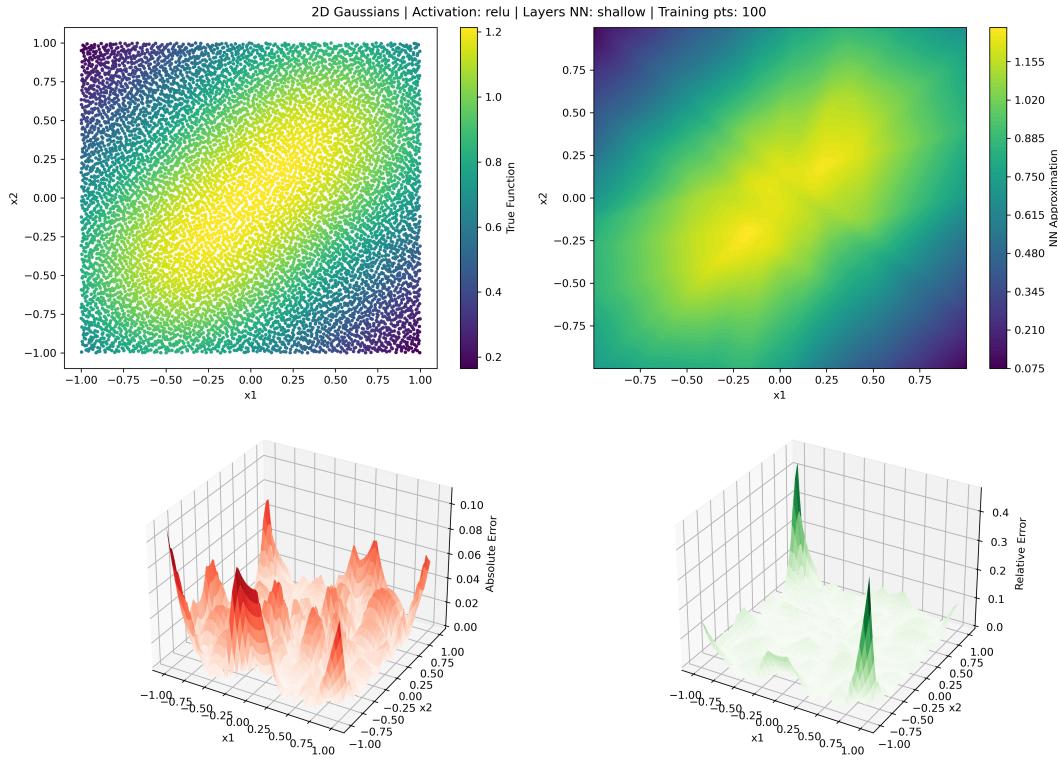


Figure 5: 2D Gaussians ($k = 1$) function results for shallow ReLU.

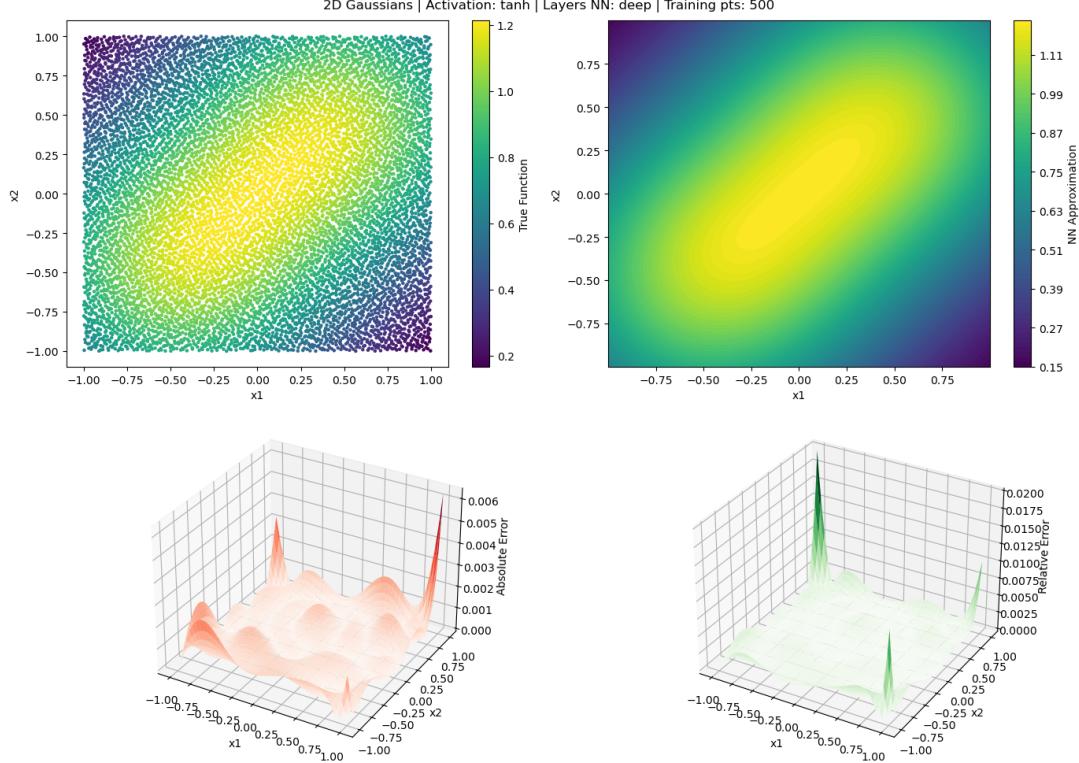


Figure 6: 2D Gaussians ($k = 1$) function results for deep tanh which achieves the lowest mean absolute error and benefits from its ability to capture multi-modal smooth structures.

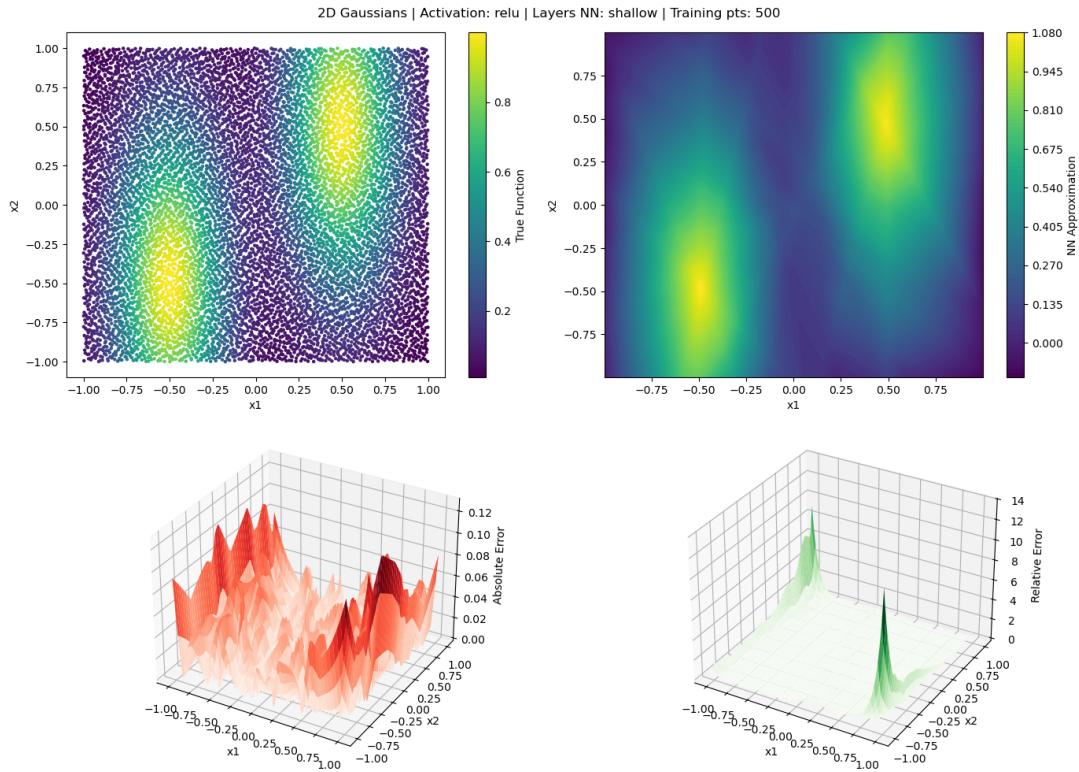


Figure 7: 2D Gaussians ($k = 10$) function results for shallow ReLU.

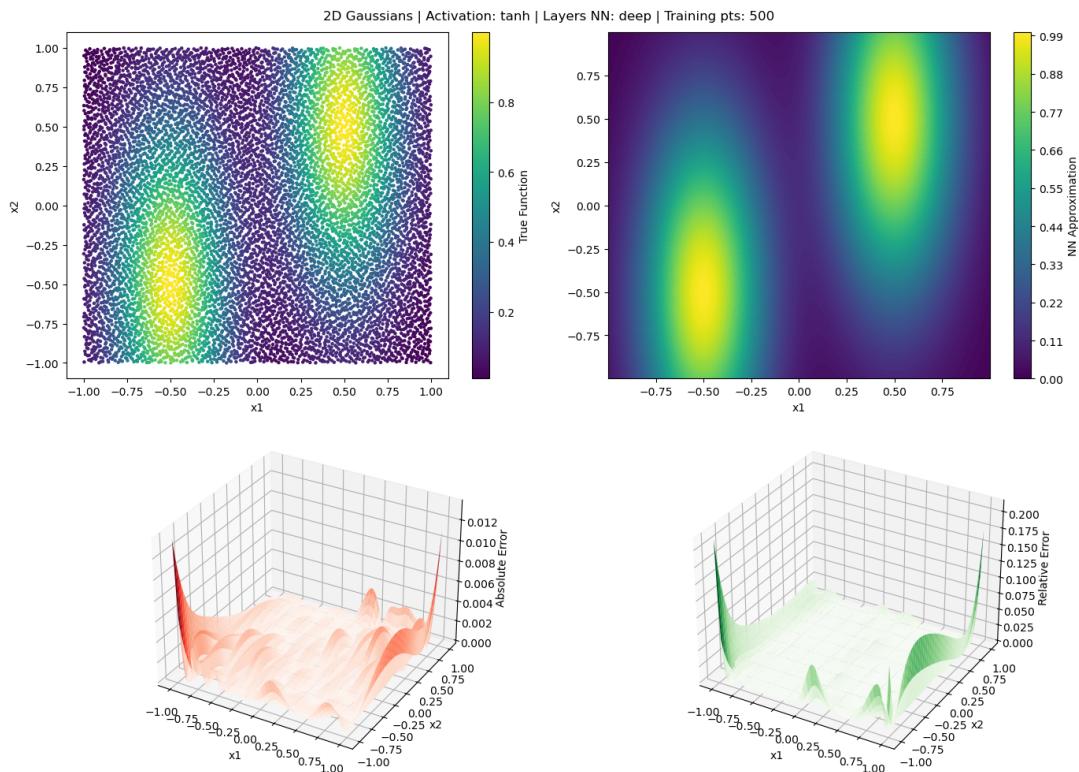


Figure 8: 2D Gaussians ($k = 10$) function results for deep tanh.

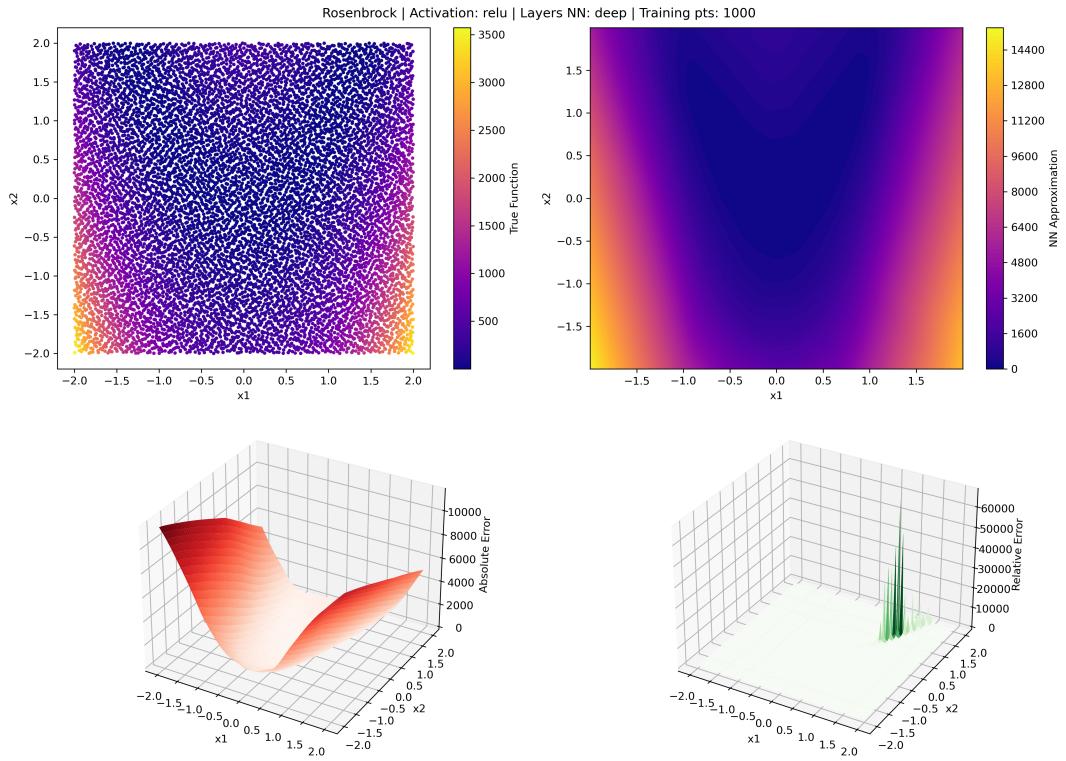


Figure 9: Rosenbrock function results for shallow ReLU which performs best and hand the complex non-linear valley effectively after y-normalization.

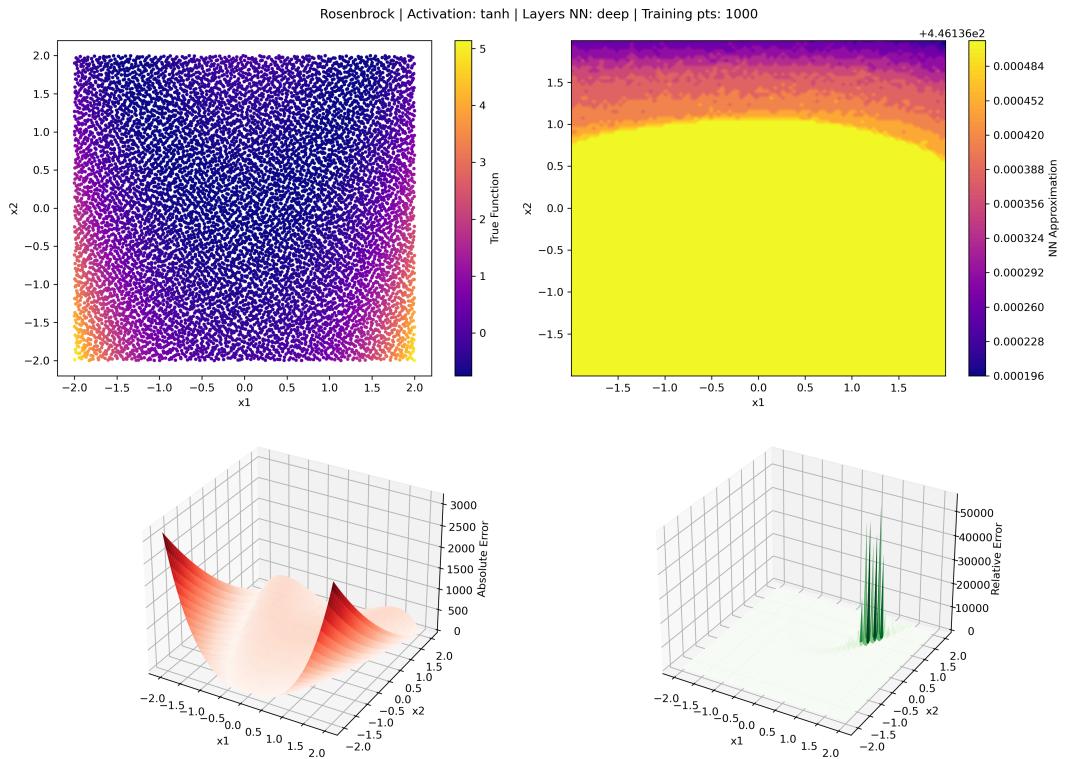


Figure 10: Rosenbrock function results for deep tanh.

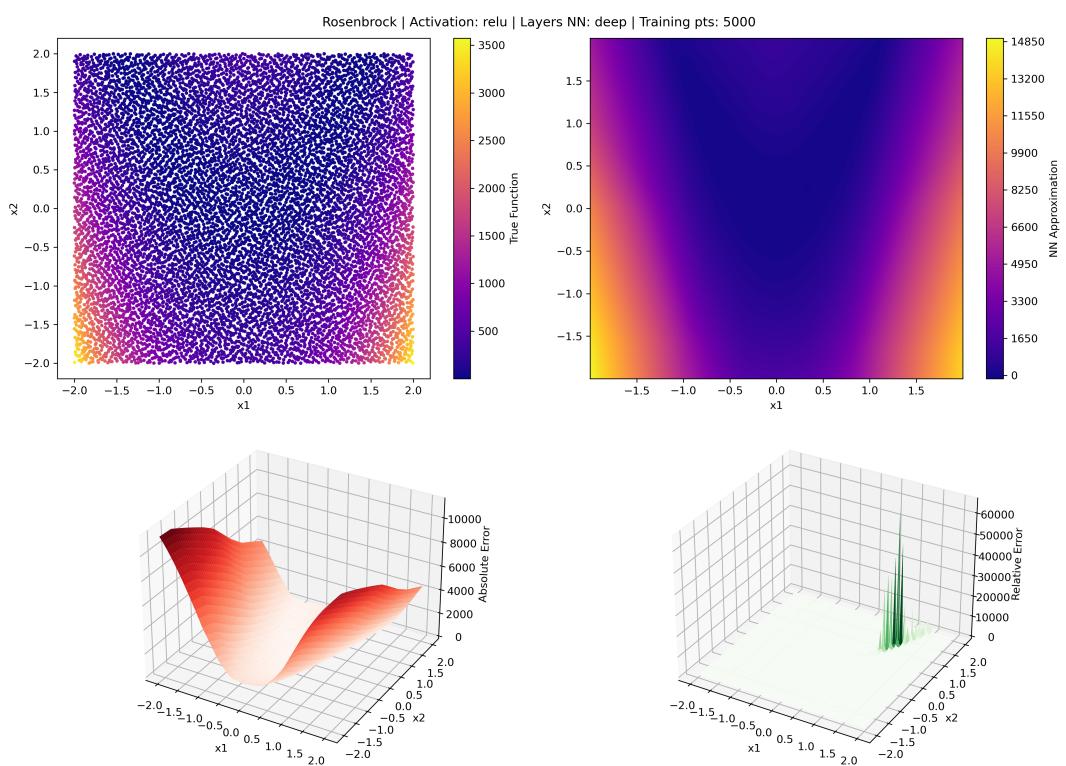


Figure 11: Rosenbrock function results for deep ReLU.